

# Name: Shishu

## Registration No.:2020CA089

### 1. PART 1: THREE PROCESSES - A PARENT AND TWO CHILDREN

In the first part you will write a small user-land program that has a parent process forking two child processes, resulting in a total of three processes: the parent process and the two children.

Put the previous code into a file named `race.c` inside your Xv6 source code folder. Add `_race\` to the `UPROGS` variable inside your Makefile. Compile and run Xv6 Run the user-land program inside Xv6 by typing `race` at the Xv6 prompt. Notice the order of execution of the three processes. Run the program multiple times.

- Do you always get the same order of execution?
  - Yes, order of execution's is same every time.
- Does Child 1 always execute (print Child 1 Executing) before Child 2?
  - No, when we add sleep statement before the child 1 print statement then first child 2 is executing.

#### Race.c

```
#include "types.h"
#include "stat.h"
#include "user.h"

//We want Child 1 to execute first, then Child 2, and finally Parent.
int main() {
    int pid = fork(); //fork the first child
    if(pid < 0) {
        printf(1, "Error forking first child.\n");
    } else if (pid == 0) {
        sleep(5);
```

```

printf(1, "Child 1 Executing\n");
} else {
pid = fork(); //fork the second child
if(pid < 0) {
printf(1, "Error forking second child.\n");
} else if(pid == 0) {
printf(1, "Child 2 Executing\n");
} else {
printf(1, "Parent Waiting\n");
int i;
for(i=0; i< 2; i++)
wait();
printf(1, "Children completed\n");
printf(1, "Parent Executing\n");
printf(1, "Parent exiting.\n");
}
}
exit();
}

```

## Makefile:

```

258 EXTRA=\
259     mkfs.c ulib.c user.h cat.c echo.c forktest.c grep.c kill.c\
260     ln.c ls.c mkdir.c rm.c stressfs.c usertests.c wc.c ps.c zombie.c\
261     printf.c myls.c umalloc.c foo.c nice.c race.c NProcess.c\
262     README dot-bochsrc *.pl toc.* runoff runoff1 runoff.list\
263     .gdbinit.tmpl gdbutil\
264

```

```
169 UPROGS=\
170     _cat\
171     _echo\
172     _forktest\
173     _grep\
174     _init\
175     _kill\
176     _ln\
177     _ls\
178     _mkdir\
179     _rm\
180     _sh\
181     _stressfs\
182     _usertests\
183     _wc\
184     _ps\
185     _myls\
186     _NProcess\
187     _zombie\
188     _foo\
189     _nice\
190     _race\
191     _niceticket\
192
```

## Output:

```
Booting from Hard Disk..xv6...
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200
init: starting sh
$ race
Child 1 Executing
PCarhenit Waitingl
d 2 Executing
Children completed
Parent Executing
Parent exiting.
$ race
Child 1 Executing
PCarheinlt Waitingd
 2 Executing
Children completed
Parent Executing
Parent exiting.
$ race
Child 1 Executing
PCahrielndt 2 EWxaeitcinug
ting
Children completed
Parent Executing
Parent exiting.
```

After adding `sleep(5);` before line before the line where Child 1 prints "Child Executing".

```

cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200
init: starting sh
$ race
PCahreintld 2 E xWaeicutitinngg

Child 1 Executing
Children completed
Parent Executing
Parent exiting.
$ race
PChairledn t2 E xWeaciting
uting
Child 1 Executing
Children completed
Parent Executing
Parent exiting.
$ race
ChilPda r2 Eexencutt inWg
aiting
Child 1 Executing
Children completed
Parent Executing
Parent exiting.

```

### What do you notice?

-I notice that firstly on xv6 terminal when we type 'race' then child 1 is created after that child 2 is created and when both the children are completing then the parent executing and then exiting.

### Can we guarantee that Child 1 always execute before Child 2?

-It is not guaranteed that child 1 always execute before child 2 because when we add a sleep statement before child 1 then first child 2 can be executed and after that child 1 Is executed.

## 2. PART 2: SPIN LOCKS

We will start by defining a spinlock that we can use in our user-land program. Xv6 already has a spinlock (see `spinlock.c`) that it uses inside its kernel and is coded in somehow a complex way to handle concurrency caused by interrupts. We don't need most of the complexity, so will write our own light-weight version of spinlocks. We will put our code inside `ulib.c`, which includes functions accessible to user-land programs.

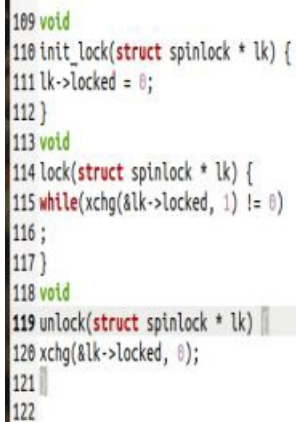
### 4. Inside `ulib.c`, add



```
Open ▾ 
1 #include "types.h"
2 #include "stat.h"
3 #include "fcntl.h"
4 #include "user.h"
5 #include "x86.h"
6 #include "spinlock.h"
7
```

to the beginning.

Also, add the following function definitions:



```
109 void
110 init_lock(struct spinlock * lk) {
111     lk->locked = 0;
112 }
113 void
114 lock(struct spinlock * lk) {
115     while(xchg(&lk->locked, 1) != 0)
116 ;
117 }
118 void
119 unlock(struct spinlock * lk) {
120     xchg(&lk->locked, 0);
121 }
122
```

We are still using the `struct spinlock` defined in `spinlock.h` but we will only use its `locked` field. Initializing the lock and unlocking it both work by setting `locked` to 0. Locking uses the

`atomic_xchg` instruction, which sets the contents of its first parameter (a memory address) to the second parameter and returns the old value of the contents of the first parameter.

5. Inside `user.h`, add the following function prototypes:

```
53 void init_lock(struct spinlock * lk);  
54 void lock(struct spinlock * lk);  
55 void unlock(struct spinlock * lk);  
56
```

to the end of the file.

Now, we have our spinlocks in place. We can use them inside `race.c`:

```
struct spinlock lk;  
init_lock(&lk);  
lock(&lk);  
//critical section  
unlock(&lk)
```

We will use condition variables to be able to make Child 2 sleep (block) until Child 1 finishes execution.

### 3. PART 3: CONDITION VARIABLES

We will use condition variables to ensure that Child 1 always executes before Child 2. We will add two system calls to Xv6: `cv_wait()` and `cv_signal()` to wait (sleep) on a condition variable and to wakeup (signal) all sleeping processes on a condition variable.

Recall that both waiting and signaling a condition variable has to be called after acquiring a lock (that's why we defined our spinlock in Part 2). `cv_wait` releases the lock before sleeping and reacquires it after waking up.

6. First, define the condition variable structure in `condvar.h` as follows.

```
#include "spinlock.h"  
struct condvar {  
    struct spinlock lk;  
};
```

A condition variable has a spin lock.

Let's then add the two system calls.

7. Inside `syscall.h`, add the following two lines:

```
32 #define SYS_cv_signal 31  
33 #define SYS_cv_wait 32
```



Inside usys.S, add:

```
41 SYSCALL(cv_signal)
42 SYSCALL(cv_wait)
...
```

Inside syscall.c, add:

```
150 [SYS_cv_wait] sys_cv_wait,
151 [SYS_cv_signal] sys_cv_signal,
...
```

And

```
[SYS_cv_wait] sys_cv_wait,
[SYS_cv_signal] sys_cv_signal,
```

Inside user.h, add

```
struct condvar;
```

to the beginning and

```
int cv_wait(struct condvar *);
int cv_signal(struct condvar *);
```

to the end of the system calls section of the file.

Our condition variable implementation depends heavily on the sleep/wakeup mechanism implemented inside Xv6 (Please read **Sleep and Wakeup** on Page 65 of the Xv6 book). We will again define a more light-weight version of the sleep function to use our light-weight spinlocks defined in Part 2 instead of



8. Inside `proc.c`, add the following function definition:

```
805 void
806 sleep1(void *chan, struct spinlock *lk)
807 {
808     struct proc *p = myproc();
809     if(p == 0)
810         panic("sleep");
811     if(lk == 0)
812         panic("sleep without lk");
813     acquire(&ptable.lock);
814     lk->locked = 0;
815     // Go to sleep.
816     p->chan = chan;
817     p->state = SLEEPING;
818     sched();
819     // Tidy up.
820     p->chan = 0;
821     release(&ptable.lock);
822     while(xchg(&lk->locked, 1) != 0)
823 ;
824 }
825
```

After a couple of sanity checks, the function acquires the process table lock `ptable.lock` to be able to call `sched()`, which works on the process table. Then, it releases the spinlock (by setting

locked to 0) and goes to sleep by setting the process state to SLEEPING, setting the channel that the process sleeps on, and switching into the scheduler by calling `sched()`. After the process wakes up, it releases the `ptable` lock and reacquires the spinlock (using the `xchg` instruction).

9. Inside `sysproc.c` add

```
#include "condvar.h"
```

to the beginning of the file and the following system call functions

```
151 int
152 sys_cv_signal(void)
153 {
154     int i;
155     struct condvar *cv;
156     argint(0, &i);
157     cv = (struct condvar *) i;
158     wakeup(cv);
159     return 0;
160 }
161 int
162 sys_cv_wait(void)
163 {
164     int i;
165     struct condvar *cv;
166     argint(0, &i);
167     cv = (struct condvar *) i;
168     sleep1(cv, &(cv->lk));
169     return 0;
170 }
```

to the end. In both functions, the code starts with retrieving the argument `(struct condvar *)` from the stack:

```
argint(0, &i);
cv = (struct condvar *) i;
```

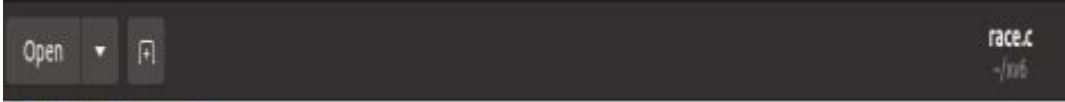
The address of the condition variable is used as the channel passed over to the `sleep1` function defined in Step 8. The address of the condition variable is unique and this is all what we need: a unique channel number to sleep and to get waked up on.

**After seeing what the two system calls do, why do you think we had to add system calls for the operations on condition variables? Why not just have these operations as functions in `ulib.c` as we did for the spinlock?**

## 4. PART 4: USING THE CONDITION VARIABLES

We can then modify race.c to use a condition variable to guarantee process ordering.

### Race.c



```
1#include "types.h"
2#include "stat.h"
3#include "user.h"
4#include "condvar.h"
5//We want Child 1 to execute first, then Child 2, and finally Parent.
6int main() {
7    struct condvar cv;
8    init_lock(&cv.lk);
9    int pid = fork(); //fork the first child
10    if(pid < 0) {
11        printf(1, "Error forking first child.\n");
12    } else if (pid == 0) {
13        sleep(5);
14        printf(1, "Child 1 Executing\n");
15        lock(&cv.lk);
16        cv_signal(&cv);
17        unlock(&cv.lk);
18    } else {
19        pid = fork(); //fork the second
20        if(pid < 0) {
21            printf(1, "Error forking second child.\n");
22        } else if (pid == 0) {
23            lock(&cv.lk);
24            cv_wait(&cv);
25            unlock(&cv.lk);
26            printf(1, "Child 2 Executing\n");
27        } else {
28            printf(1, "Parent Waiting\n");
29            int i;
30            for(i=0; i< 2; i++)
31                wait();
32            printf(1, "Children completed\n");
33            printf(1, "Parent Executing\n");
34            printf(1, "Parent exiting.\n");
35        }
36    }
37    exit();
38 }
```

Note the highlighted parts. A condition variable is declared. Its spinlock is initialized. Then Child 1

signals the condition variable after acquiring the spinlock. Child 2 sleeps on the condition variable after acquiring the spinlock as well.

Compile and run the modified race program.

- Does Child 1 always execute before Child 2?

## 5. PART 5: LOST WAKEUPS

Does it happen that the program gets stuck? This is called a **deadlock**. If Child 2 gets to sleep **after** Child 1 signals, the wakeup signal is lost (i.e., never received by Child 2). In this case, Child 2 has no way of being awaked.

To solve this problem, we need to enclose the `cv_wait()` call inside a while loop. We need some form of a flag that gets set by Child 1 when it is done executing. Child 2 will then do

```
while(flag is not set)
```

```
    cv_wait();
```

This way, even if Child 1 sets the flag and signals before Child 2 executes the while loop, Child 2 will not avoid going to sleep because the flag will be set.

The flag has to be **shared** between the two processes. We will use a **file** for that. Other methods for sharing are shared memory and pipes.

To create a file,

```
int fd = open("flag", O_RDWR | O_CREATE);
```

To write into the file,

```
write(fd, "done", 4);
```

Checking the flag has to be non-blocking. The `read` system call is blocking. Reading the size of the file is not. So, we will check the flag by reading the file size. To read the size of a file,

Now, we are ready to write the while loop inside Child 2. It will loop until the file size is greater than zero, which happens when Child 1 writes "done" into the file after it finishes execution.

The new `race.c` is:

```

#include "user.h"
#include "condvar.h"
#include "fcntl.h"

//We want Child 1 to execute first, then Child 2, and finally Parent.
int main() {
    struct condvar cv;
    int fd = open("flag", O_RDWR | O_CREATE);
    init_lock(&cv.lk);
    int pid = fork(); //fork the first child
    if(pid < 0) {
        printf(1, "Error forking first child.\n");
    } else if (pid == 0) {
        sleep(5);
        printf(1, "Child 1 Executing\n");
        lock(&cv.lk);
        write(fd, "done", 4);
        cv_signal(&cv);
        unlock(&cv.lk);
    } else {
        pid = fork(); //fork the second
        if(pid < 0) {
            printf(1, "Error forking second child.\n");
        } else if(pid == 0) {
            lock(&cv.lk);
            struct stat stats;
            fstat(fd, &stats);
            printf(1, "file size = %d\n", stats.size);
            while(stats.size <= 0){
                cv_wait(&cv);
                fstat(fd, &stats);
                printf(1, "file size = %d\n", stats.size);
            }
            unlock(&cv.lk);
            printf(1, "Child 2 Executing\n");
        } else {
            printf(1, "Parent Waiting\n");
            int i;
            for(i=0; i< 2; i++)
                wait();
            printf(1, "Children completed\n");
            printf(1, "Parent Executing\n");
            printf(1, "Parent exiting.\n");
        }
    }
    close(fd);
    unlink("flag");
}

```

Note the highlighted parts and also note that we are closing the file and deleting it before the parent exits. This is to start afresh the next time we run the program.

Compile and run `race.c` many times.

- Is it always the case that Child 1 executes before Child 2?
- Do you observe deadlocks?

Of course, synchronization bugs cannot be ruled out by running a program many times. Formal proof is typically the preferred way especially for safety- and mission-critical systems. There are tools that help with this kind of formal proofs.

## 6. SUBMISSION INSTRUCTIONS

Follow the steps of this lab and submit all the files that you had to modify in a zipped archive by uploading the zipped file to Lab 2 submission page on CourseWeb. The deadline is listed on the Lab 2 page on CourseWeb.